Full technology stack

Layer	Tool / Library	Why chosen
Embedding	all-MiniLM-L6-v2	Offline, fast, 80MB, great semantic similarity
Vector search	FAISS (Facebook)	In-memory, no server, billion-scale fast search
LLM (generator)	Mistral 7B / LLaMA 3	Open-source, runs offline, strong generation quality
NLI / entailment	DeBERTa NLI cross-encoder	Best entailment accuracy, lightweight, offline
Semantic score	BERTScore	Token-level overlap, no external calls needed
Orchestration	LangChain / custom Python	Modular, easy pipeline building, well documented